

# Lecture 10


# Deadlock

Prof. Yinqian Zhang

Spring 2025

# Starvation vs. Deadlock

- Starvation vs. Deadlock
  - Starvation: thread waits indefinitely
    - Low-priority thread waiting for resources constantly in use by high-priority threads
  - Deadlock: circular waiting for resources
    - Thread A owns Res 1 and is waiting for Res 2
    - Thread B owns Res 2 and is waiting for Res 1
  - Deadlock  $\Rightarrow$  Starvation but not vice versa
    - Starvation can end (but does not have to)
    - Deadlock cannot end without external intervention



# Conditions for Deadlock

- Deadlock will not always happen
  - Need the exactly right timing
  - Bugs may not exhibit during testing
- Deadlocks occur with multiple resources
  - Cannot solve deadlock for each resource independently
  - System with 2 disk drives and two threads
  - Each thread needs 2 disk drives to function
  - Each thread gets one disk and waits for another one

## Process A

sem\_wait(**x**)  
sem\_wait(**y**)

sem\_post(**y**)  
sem\_post(**x**)

## Process B

sem\_wait(**y**)  
sem\_wait(**x**)


sem\_post(**x**)  
sem\_post(**y**)

# Four Requirements for Deadlock

- Mutual exclusion
  - Only one thread at a time can use a resource.
- Hold and wait
  - Thread holding at least one resource is waiting to acquire additional resources held by other threads
- No preemption
  - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
- Circular wait
  - There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
    - $T_1$  is waiting for a resource that is held by  $T_2$
    - $T_2$  is waiting for a resource that is held by  $T_3$
    - ...
    - $T_n$  is waiting for a resource that is held by  $T_1$

# Resource-Allocation Graph


- System Model
  - A set of Threads  $T_1, T_2, \dots, T_n$
  - Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
  - Each resource type  $R_i$  has  $W_i$  instances
  - Each thread utilizes a resource as follows:
    - Request() / Use() / Release()
- Resource-Allocation Graph:
  - $V$  is partitioned into two types:
    - $T = \{T_1, T_2, \dots, T_n\}$ , the set threads in the system.
    - $R = \{R_1, R_2, \dots, R_m\}$ , the set of resource types in system
  - request edge – directed edge  $T_i \rightarrow R_j$
  - assignment edge – directed edge  $R_j \rightarrow T_i$




# Resource Allocation Graph Examples

- Recall:


- request edge - directed edge  $T_i \rightarrow R_j$
- assignment edge - directed edge  $R_j \rightarrow T_i$



Simple Resource Allocation Graph



Allocation Graph With Deadlock




Allocation Graph with Cycle, but No Deadlock

# Methods for Handling Deadlocks

- Allow system to enter deadlock and then recover
  - **Deadlock detection**
  - Requires deadlock detection algorithm
  - Some technique for forcibly preempting resources and/or terminating tasks
- Ensure that system will never enter a deadlock
  - **Deadlock prevention**
  - Need to monitor all resource acquisitions
  - Selectively deny those that might lead to deadlock
- Ignore the problem and pretend that deadlocks never occur in the system
  - Used by most operating systems, including UNIX

# Deadlock Detection with Resource Allocation Graphs

- Only one of each type of resource  $\Rightarrow$  look for cycles
- More than one resource of each type
  - More complex deadlock detection algorithm
  - Next page





# Several Instances per Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $\text{Request}[i_j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively Initialize:
  - (a) **Work** = **Available**
  - (b) For  $i = 1, 2, \dots, n$ , if **Allocation<sub>i</sub>**  $\neq 0$ , then  
**Finish**[ $i$ ] = false; otherwise, **Finish**[ $i$ ] = true
2. Find an index  $i$  such that both:
  - (a) **Finish**[ $i$ ] == false
  - (b) **Request<sub>i</sub>**  $\leq$  **Work**

If no such  $i$  exists, go to step 4
3. **Work** = **Work** + **Allocation<sub>i</sub>**;  
**Finish**[ $i$ ] = true  
go to step 2
4. If **Finish**[ $i$ ] == false, for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if **Finish**[ $i$ ] == false, then  $P_i$  is deadlocked

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $\text{Finish}[i] = \text{true}$  for all  $i$

## Example (Cont.)

- $P_2$  requests an additional instance of type C

	<u>Request</u>		
	A	B	C
$P_0$	0	0	0
$P_1$	2	0	2
$P_2$	0	0	1
$P_3$	1	0	0
$P_4$	0	0	2

- State of system?
  - Can reclaim resources held by process  $P_0$  (not deadlocked), but insufficient resources to fulfill other processes; requests
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$

# What if Deadlock Detected?

- Terminate process, force it to give up resources
  - Shoot a dining philosopher !?
  - But, not always possible
- Preempt resources without killing off process
  - Take away resources from process temporarily
  - Does not always fit with semantics of computation
- Roll back actions of deadlocked process
  - Common technique in databases (transactions)
  - Of course, deadlock may happen once again

# Deadlock Prevention

- Try to ensure at least one of the conditions cannot hold to prevent deadlock
  - Remove “Mutual Exclusion”: not possible for non-sharable resources
  - Remove “Hold and Wait” – must guarantee that whenever a process requests a resource, it does not hold any other resources
    - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none
  - Low resource utilization; starvation possible

# Deadlock Prevention

- Try to ensure at least one of the conditions cannot hold to prevent deadlock
  - Remove “No preemption”
    - If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
    - Preempted resources are added to the list of resources for which the process is waiting
    - Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

# Deadlock Prevention

- Try to ensure at least one of the conditions cannot hold to prevent deadlock
  - Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration
  - $R = \{R_1, R_2, \dots, R_m\}$
  - One to one function  $F: R \rightarrow N$
  - If a process request a resource  $R_i$ , it can request another resource  $R_j$  if and only if  $F(R_i) < F(R_j)$
  - Or, it must first release all resource  $R_i$  such that  $F(R_i) \geq F(R_j)$



# Deadlock Avoidance


- Requires that the system has some additional a priori information available
  - Simplest and most useful model requires that each process declare the maximum number of resources of each type that it may need
  - The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a **circular-wait condition**
  - Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes

# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that **for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$** 
  - If what  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

# Safe, Unsafe, Deadlock State

- If a system is in safe state => no circular wait => no deadlocks
- If a system is in unsafe state => possibility of deadlock
- Deadlock avoidance => ensure that a system will never enter an unsafe state.



# Banker's Algorithm

- Multiple instances of each resource type
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

# Banker's Algorithm (Cont'd)

- Let  $n$  = number of processes, and  $m$  = number of resources types.
- **Available:** Vector of length  $m$ . If **available**  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If **Max**  $[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If **Allocation** $[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If **Need** $[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$\text{Need} [i,j] = \text{Max}[i,j] - \text{Allocation} [i,j]$$

# Banker's Algorithm: Safety Algorithm

1. Let **Work** and **Finish** be vectors of length m and n, respectively.

Initialize:

**Work** = **Available**

**Finish** [i] = false for i = 0, 1, ..., n- 1

2. Find an index i such that both:

(a) **Finish** [i] = false

(b) **Need**<sub>i</sub> ≤ **Work** (i.e., for all k, **Need**<sub>i</sub>[k] ≤ **Work**[k])

If no such i exists, go to step 4

3. **Work** = **Work** + **Allocation**<sub>i</sub>

**Finish**[i] = true

go to step 2

4. If **Finish** [i] == true for all i, then the system is in a safe state

# Resource-Request Algorithm for Process $P_i$

**Request** = request vector for process  $P_i$ . If **Request** <sub>$i$</sub>  [ $j$ ] =  $k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If **Request** <sub>$i$</sub>   $\leq$  **Need** <sub>$i$</sub>  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If **Request** <sub>$i$</sub>   $\leq$  **Available**, go to step 3. Otherwise,  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

**Available** = **Available** - **Request**;

**Allocation** <sub>$i$</sub>  = **Allocation** <sub>$i$</sub>  + **Request** <sub>$i$</sub> ;

**Need** <sub>$i$</sub>  = **Need** <sub>$i$</sub>  - **Request** <sub>$i$</sub> ;

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>MAX</u>			<u>Available</u>		
	$A$	$B$	$C$	$A$	$B$	$C$	$A$	$B$	$C$
$P_0$	0	1	0	7	5	3	3	3	2
$P_1$	2	0	0	3	2	2			
$P_2$	3	0	2	9	0	2			
$P_3$	2	1	1	2	2	2			
$P_4$	0	0	2	4	3	3			



## Example (Cont'd)

- The content of the matrix Need is defined to be **Max - Allocation**

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	3 3 2
$P_1$	2 0 0	1 2 2	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria

## Example: $P_1$ Request (1,0,2)

- Check that **Request**  $\leq$  **Available**, that is,  $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	1	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can request for (3,3,0) by  $P_4$  be granted?
- Can request for (0,2,0) by  $P_0$  be granted?

**Thank you!**

